

Introduction:

- When not to use SQL:
- Needing high availability in the data: system always up; no SFT
 - Have large amounts of data
 - Need linear scalability: vertical scaling → adding more nodes, linearly increases performance.
 - Non ACIDity
 - ↳ Need fast read & writes

NoSQL Databases Explained:

↳ Not Only; ouputs unstructured data

Key-Value Database: uses key associated with one collection of values.

↳ may be a book

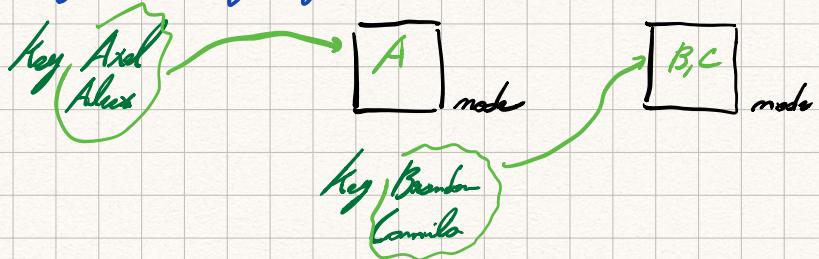
↳ stored in a DB

A key-value → has no = query language "

Store, retrieve & update data → set, get, delete

A key-value → "key" selection is critical → queries are performed just in box to it.

→ Similar "keys" go to nearby physical location → can lead to failure if no redundancy



A Consistency issue → if two machines have the key, there are problems of value consistency.

Column-family Databases: stores data in column families or tables. These tables have many columns associated with a particular row.

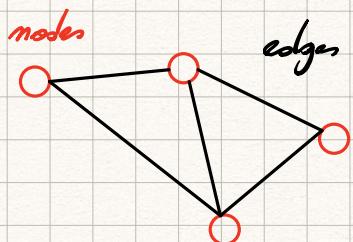
Row Key X	Column 1 name1.value1	Column 2 name2.value2	Column N nameN.valueN
-----------	--------------------------	--------------------------	--------------------------

Row Key Y	Column 1 name1.value1	Column 2 name2.value2	Column N nameN.valueN
-----------	--------------------------	--------------------------	--------------------------

a compression: stores data efficiently through data compression and by using data partitioning

- Agg. Queries due to structure; aggregation operations are performant
- Scalability: more scalable than other DB.
- Fast to read & query.

Graph DB: store data in the form of nodes & edges



Nodes → Entities

Edges → Relationships

- directional significance
- properties

- ⚠ → Adding New node: easy ⚡
- Modifying current node relationships. hard 🚫

Choosing the right = NoSQL Database:

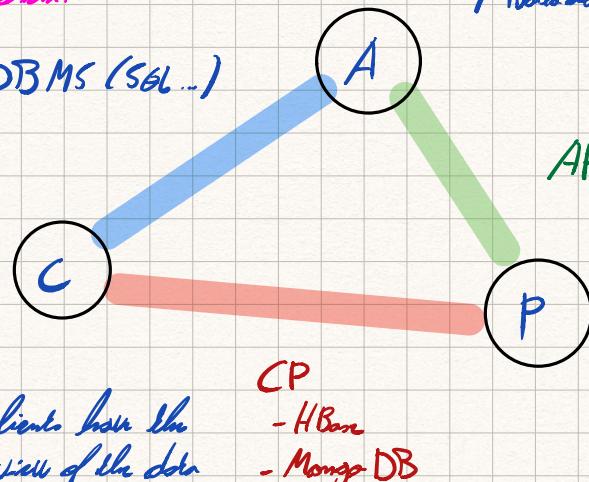
→ Data you need to store

→ Data model

→ Data reading expectations

CAP Theorem:

CA: RDBMS (SQL...)



Availability: each client can always read & write

AP: Consistency

P Partition Tolerant: system works well despite physical network partition

Consistency: all clients have the same view of the data

Distributed Database:

→ In order to have higher availability → You need to have copies of your data.

All copies may not be up-to-date



Eventual consistency: if no updates to a data item, eventually (i.e. not immediately), all access to that item return last updated value

In a consistency model in distributed computing

↳ If no new updates are made to new data, eventually all access will return last updated value.



When the system achieves eventual consistency, the system has converged / replica convergence.

→ Also called optimistic replication



Conflict Resolution: systems of db reconciliation

↳ anti-entropy: exchanging versions or updates

↳ reconciliation: choosing an appropriate first disk when concurrent updates have occurred.

Strategies →

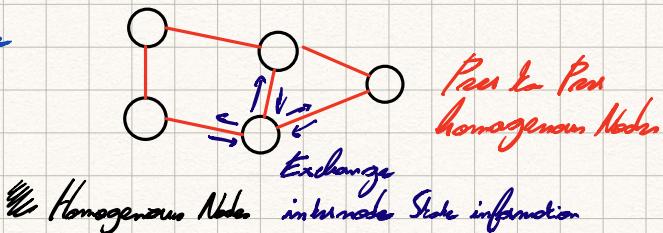
- "last writer wins"
- "first writer wins"

When

→ read repair: correction is done when a read finds an inconsistency
write repair: correction takes place during write operation
asynchronous repair: correction is not part of read and write.

Apache Cassandra

→ Structure



- Data: a partitioned row store database
 - Requires a primary key

→ When a partition arrives: Written in (GEL / Similar to SGT),

the resulting node acts as "coordinator".

DATA

1. Sequential writer commit
by copy-on-write activity

↓
FLOW

2. Data is then indexed and written in a memory structure
↳ memtable

↓ When full

3. Data written to a disk in
SSTable

↳ Automatically → Replicated
↳ Partitioned

→ Data Replication: When an old copy of data is detected, while reading, it returns the most updated value and performs a "read repair" (slows reading speed to enforce consistency).

→ Keyspace: × Replication factor: n° of machines with a copy of data.

× Replication Strategy: strategy for replica placement

↳ Simple strategy
↳ All Network Topology Strategy
↳ Network Topology Strategy

× Column families: a "table", each one containing many rows.

↳ Not defined from creation; columns can be added at any moment, to a give column family.

Column:

name: byte[] ; value: byte[] ; clock: byte[]

Super Column:

name: byte[], cols: map<byte[], column>



↳ "Supercolumn" column are physically saved together, blocks combine column that are usually queried together.

CAP Theorem:

Theorem in computer science that states it is impossible for a distributed data storage to simultaneously offer more than two out of the 3 guarantees:

- Consistency: every read from the database gets the latest (and correct) piece of data or an error.
- Availability: every request is received and responded; no data (most recent) guarantee.
- Partition Tolerance: the system continues to work regardless of losing network connectivity between nodes.

CAP Consistency vs ACID RDB?

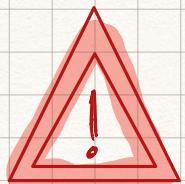
ACID relates to database rules: Atomic all components of a transaction are treated together either all done or consistent transaction must follow DB rules (schema constraints)
Isolated concurrent execution of transactions is translated into a linear order
Durable committed transactions are NOT UNDONE

CAP: relates to the state of the system "vis-a-vis" replicas of a same value, across the node cluster for a given specific time.

Denormalizing in Apache Cassandra:

- Denormalization is critical in Apache Cassandra there are **NO JOINS**.

↳ Need to think about your queries in order to build the tables.



Normalized Tables \Rightarrow Denormalized Tables
RDB

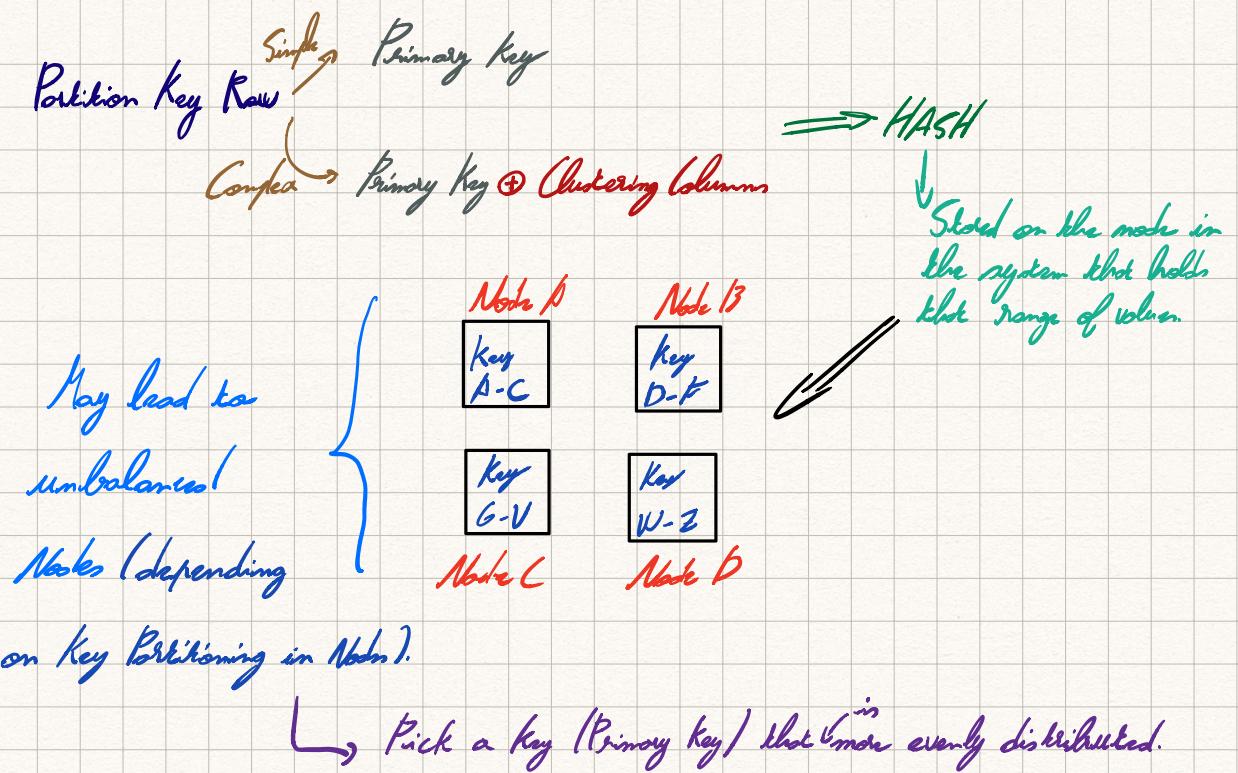
@@ Which are your queries?
=

Data modeling:

- Denormalizing is not just ok - **MUST**
- Denormalizing done for fast reads
- It is optimized for fast writes (reform read/repair).
- **ALWAYS** think first
- One Table for query \rightarrow good strategy
- **NO JOINS**

PRIMARY KEY:

- Must be unique
- Composed of either just the PARTITION KEY // PARTITION KEY Ⓛ CLUSTERING Cols.
- Simple "Primary Key" → Just the PARTITION KEY
- PARTITION KEY will determine distribution of data across the system.



Clustering Columns:

- Clustering column will sort the data in sorted ascending order.
- More than one clustering column can be added
- Clustering column will sort in order of how they were added to the primary key.

⚠️ Can use as many clustering columns as required; use them in the "SELECT STATEMENT" on the same order

Example:

CREATE TABLE music-library (year int, artist-name text, album-name text, ...,
PRIMARY KEY ((year), artist-name, album-name))

Partition key

↳ Clustering column (ordered ascending order).

The WHERE clause:

a) Data modeling in Apache Cassandra is query focused!

The Partition key must be used; any "clustering column" must be used in the order employed in PARTITION KEY definition.